

# From Scratch to Smart

A 3-Phase Demo: Pre-training, SFT, and RLHF

# The "Baby" Model: Our Starting Point

“ We start with a ‘ToyLLM’. Its ‘brain’ (weights) is completely random. When we ask it a question, its answer is meaningless garbage. ”

Prompt: 'alice followed'  
Response: 'alice rabbit queen cat curious...'

# The Core NLP Building Blocks



## Tokenization

A "vocabulary" to turn words like 'alice' into simple numbers like [12].



## Embeddings

A "dictionary" that maps these numbers to \*meaning vectors\* that the model can learn.

# Phase 1: Foundational Pre-training

## Goal: Learn Language Structure

We teach the model the basics of language by making it predict the next word over and over.

- **Method:** Next-Word Prediction
- **Data:** Simple sentences from "Alice in Wonderland."
- **Example:** Show it "alice followed the..." and teach it to predict "...white".



# Pre-training Result: A "Completer"

“ The model is no longer random! It completes sentences based on the 'Alice' facts it learned.  
But it can't \*answer\* questions. ”

Prompt: 'who did alice...'  
Response: 'who did alice follow the...'

# Phase 2: Supervised Fine-Tuning (SFT)

## Goal: Teach Factual Answers

We teach the model to be a helpful "expert" by showing it "flashcards" of (Question, Answer) pairs.

- **Method:** Instruction Following
- **Data:** Q&A pairs.
- **Example:**  
Prompt: 'who did alice follow'  
Answer: 'alice followed the white rabbit'



# SFT Result: A "By-the-Book" Expert

## Good: Trained Questions

It perfectly answers the \*specific questions\* it was trained on.

Prompt: 'what did the cat do'

Response: 'the cheshire cat smiled'

## Bad: New Questions

When asked a \*new\* question, it falls back on its pre-trained "completer" knowledge.

Prompt: 'who was alice'

Response: 'alice was very curious'

# Phase 3: Reinforcement Learning (RLHF)

## Goal: Teach Preferences

We teach the model what a "good" answer *\*feels\** like, even for new questions.

1. Model generates an answer.
2. A "Reward Model" (Judge) scores it.
3. We update the model to get more "points" from the judge in the future.



# RLHF Result: An "Aligned" Model

“  
The model is now "steered" by our preferences. Our judge "rewarded" the answer '**alice was very curious**' , teaching the model that this is a good, helpful response. ”

Prompt: 'who was alice'  
Response: 'alice was very curious'

# The Model's Evolution: A Comparison

## Pre-Trained

**Prompt:** 'who did alice follow'

**Response:** 'who did alice follow the...'

**Prompt:** 'who was alice'

**Response:** 'who was alice said the...'

## SFT

**Prompt:** 'who did alice follow'

**Response:** 'alice followed the white rabbit'

**Prompt:** 'who was alice'

**Response:** 'alice was very curious'

## RLHF

**Prompt:** 'who did alice follow'

**Response:** 'alice followed the white rabbit'

**Prompt:** 'who was alice'

**Response:** 'alice was very curious'

# Key Demo Takeaways

- 💡 **Pre-training** teaches \*language\* and \*facts\*. It creates a "completer."
- 🎓 **SFT** teaches \*instruction following\*. It creates an "expert."
- 👍 **RLHF** teaches \*preferences\*. It creates an "aligned" assistant.

# Questions?

# Image Sources



<https://media.istockphoto.com/id/1412637153/photo/magic-book-with-open-pages-and-abstract-lights-shining-in-darkness-literature-and-fairy-tale.jpg?s=612x612&w=0&k=20&c=rc3Bh74ePX8QPGnOi5oeEYowFaJYsseXu8A0dGFqRzl=>

Source: [www.istockphoto.com](http://www.istockphoto.com)

---



[https://png.pngtree.com/png-clipart/20250422/original/pngtree-stylized-q-a-speech-bubbles-illustration-png-image\\_20849901.png](https://png.pngtree.com/png-clipart/20250422/original/pngtree-stylized-q-a-speech-bubbles-illustration-png-image_20849901.png)

Source: [pngtree.com](http://pngtree.com)

---



<https://ecdn.teacherspayteachers.com/thumbitem/Thumbs-Check-for-Understanding-1609884359/original-268414-1.jpg>

Source: [www.teacherspayteachers.com](http://www.teacherspayteachers.com)